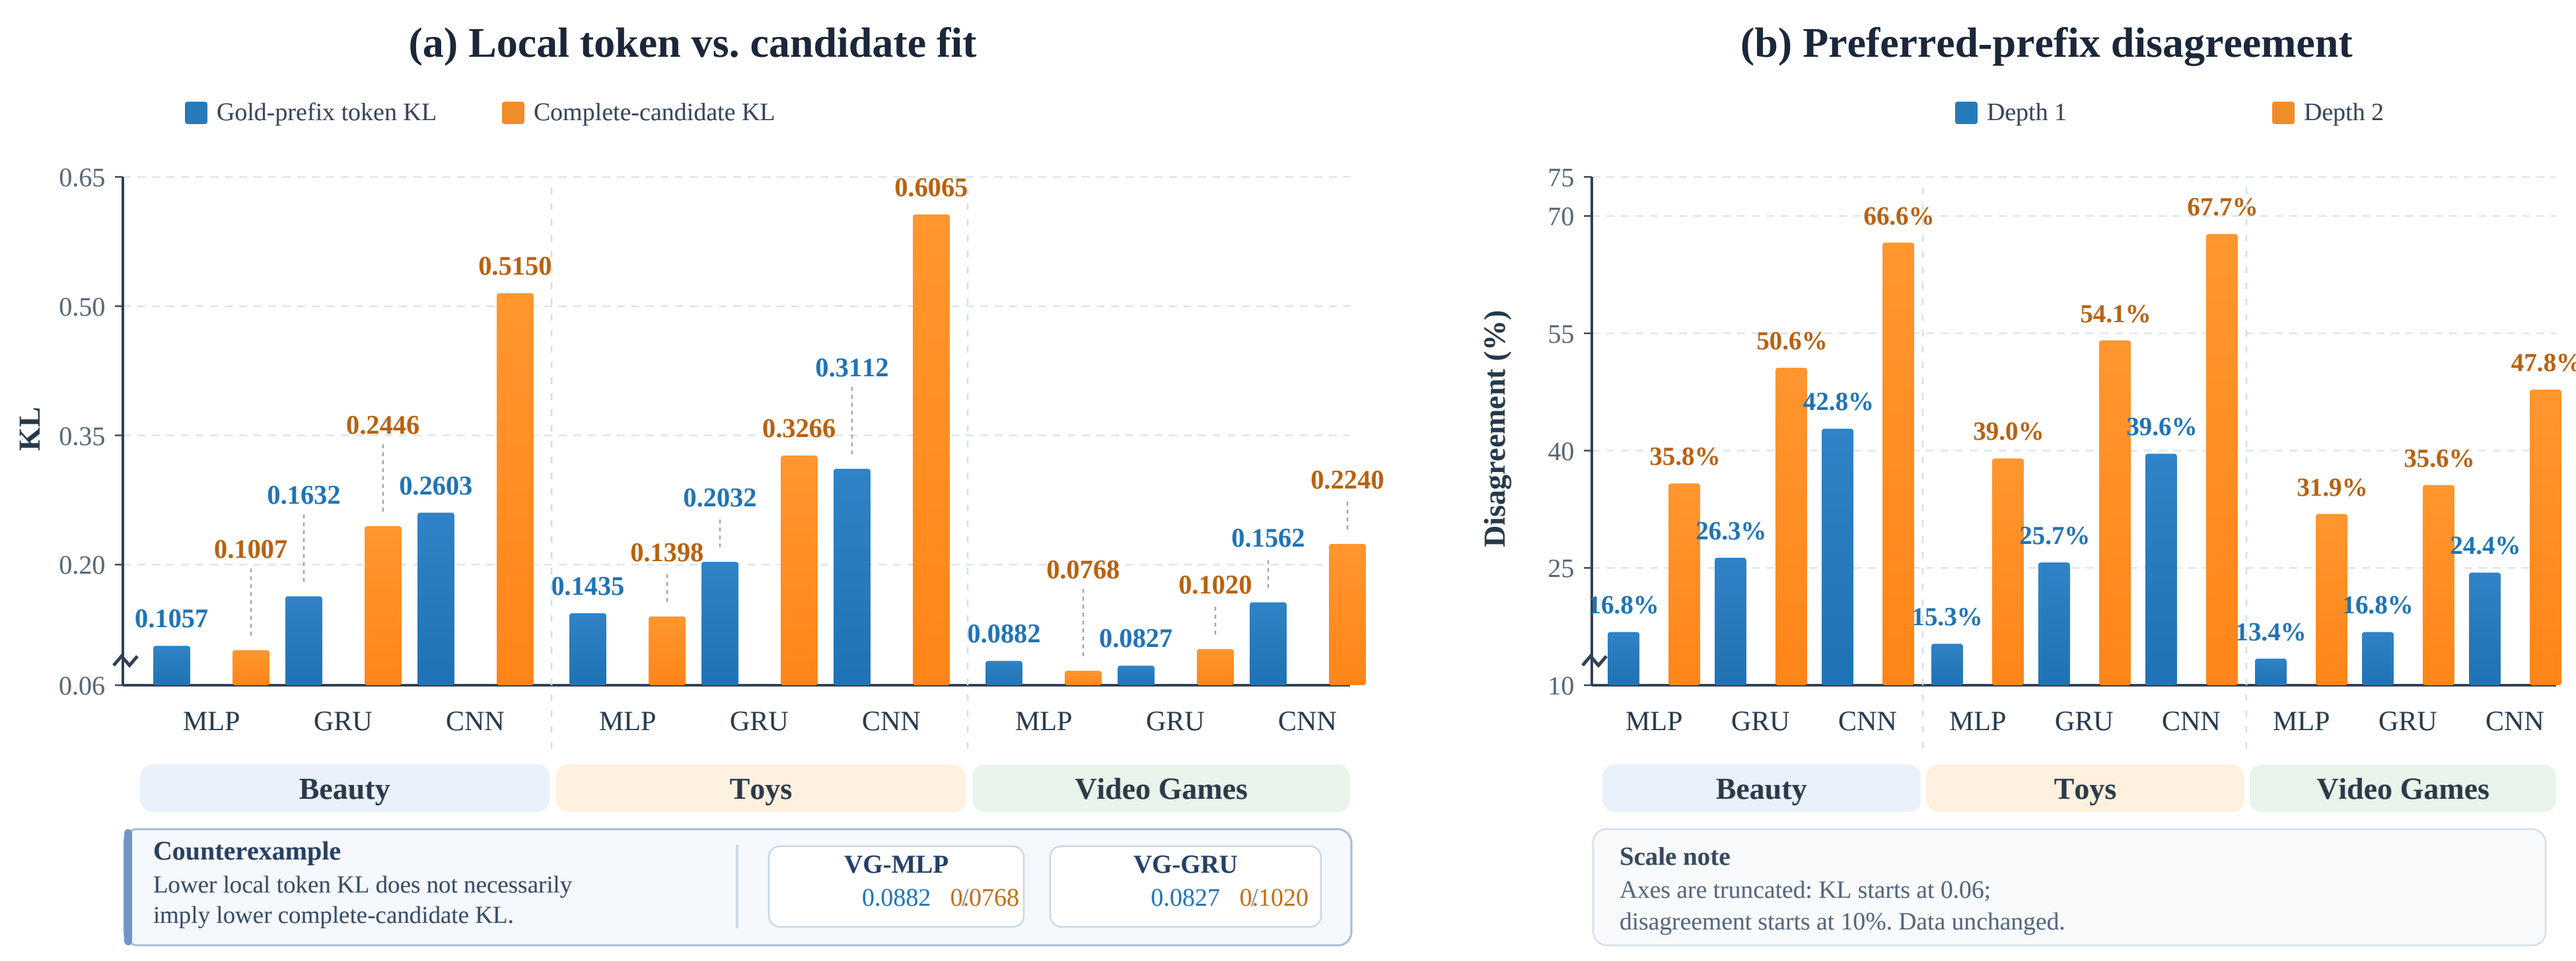
(a) Local token vs. candidate fit
Gold-prefix token KL
Complete-candidate KL
KL
0.65
0.50
0.35
0.20
0.06
0.1057
0.1007
0.1632
0.2446
0.2603
0.5150
0.1435
0.1398
0.2032
0.3266
0.3112
0.6065
0.0882
0.0768
0.0827
0.1020
0.1562
0.2240
MLP
GRU
CNN
Beauty
Toys
Video Games
Counterexample
Lower local token KL does not necessarily imply lower complete-candidate KL.
VG-MLP
0.0882 0.0768
VG-GRU
0.0827 0.1020
(b) Preferred-prefix disagreement
Depth 1
Depth 2
Disagreement (%)
75
70
55
40
25
10
16.8%
35.8%
26.3%
50.6%
42.8%
66.6%
15.3%
39.0%
25.7%
54.1%
39.6%
67.7%
13.4%
31.9%
16.8%
35.6%
24.4%
47.8%
MLP
GRU
CNN
Beauty
Toys
Video Games
Scale note
Axes are truncated: KL starts at 0.06; disagreement starts at 10%. Data unchanged.